

Training Neural Networks 1

Katharina Kann — CSCI/LING5832

Let's Talk about Weights...

- So... where do weights come from?

Let's Talk about Weights...

- So... where do weights come from?
- We'll learn the weights from a training set labeled with the right answer.
 - For instance, for sentiment analysis, we can use 1 and 0 as labels to stand for the right answer

Let's Talk about Weights...

- So... where do weights come from?
- We'll learn the weights from a training set labeled with the right answer.
 - For instance, for sentiment analysis, we can use 1 and 0 as labels to stand for the right answer
- The system answers will be between 0 and 1
 - (Why?)

Loss Functions

Let's Talk about Weights...

- To learn the weights, we need some measure of how well (or badly) we're doing with a current set of weights.
 - We'll call that measure a **loss function**

Let's Talk about Weights...

- To learn the weights, we need some measure of how well (or badly) we're doing with a current set of weights.
 - We'll call that measure a **loss function**
- The lower the loss the better we're doing

Let's Talk about Weights...

- To learn the weights, we need some measure of how well (or badly) we're doing with a current set of weights.
 - We'll call that measure a **loss function**
- The lower the loss the better we're doing
- We want to minimize the loss

Loss Functions

- There are lots of ways to measure how well we're doing on a test/validation set.
 - How many examples are right?

Loss Functions

- There are lots of ways to measure how well we're doing on a test/validation set.
 - How many examples are right?
 - How close our answers are to the correct answers?
 - What does close mean?

Cross-Entropy Loss

- For a given example, we'll go with the probability assigned to the right answer by the model

Cross-Entropy Loss

- For a given example, we'll go with the probability assigned to the right answer by the model
- In the binary case, the right answer is always either 1 or 0.
 - The model produces a number between 1 and 0, that's the starting point for the loss.

Cross-Entropy Loss

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- y is the correct answer
- \hat{y} is the system answer

Cross-Entropy Loss

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- If the correct answer for an example is 1 and the system output is .7 then the probability of the correct answer is...

Cross-Entropy Loss

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- If the correct answer for an example is 1 and the system output is .7 then the probability of the correct answer is...
 - .7

Cross-Entropy Loss

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- If the correct answer for an example is 0 and the system output is .8 then the probability of the correct answer is...

Cross-Entropy Loss

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- If the correct answer for an example is 0 and the system output is .8 then the probability of the correct answer is...
 - .2


Logs

- Probabilities aren't exactly what we want. We want worse performance to have high loss and good performance to have low loss.


Logs

- Probabilities aren't exactly what we want. We want worse performance to have high loss and good performance to have low loss.
- So we'll take the negative of the log of the probability assigned to the correct answer as the loss.

Log Probability



Log Probability



Cross-Entropy Loss for Logistic Regression

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

Cross-Entropy Loss for Logistic Regression

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

$$\begin{aligned}\log p(y|x) &= \log [\hat{y}^y (1 - \hat{y})^{1-y}] \\ &= y \log \hat{y} + (1 - y) \log(1 - \hat{y})\end{aligned}$$

Cross-Entropy Loss for Logistic Regression

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

$$\begin{aligned}\log p(y|x) &= \log [\hat{y}^y (1 - \hat{y})^{1-y}] \\ &= y \log \hat{y} + (1 - y) \log(1 - \hat{y})\end{aligned}$$

$$L_{CE}(w, b) = -[y \log \sigma(w \cdot x + b) + (1 - y) \log(1 - \sigma(w \cdot x + b))]$$

In-Class Exercise

- You are classifying tweets into positive and negative, using a logistic regression classifier.
- Compute the cross-entropy loss for
 - $y = 1, \hat{y} = 0.66$
 - $y = 1, \hat{y} = 0.7$
 - $y = 0, \hat{y} = 0.2$
 - $y = 0, \hat{y} = 0.8$
- How many examples have been predicted correctly by your classifier?

Learning

- We want to find the weights that minimize the average loss of an entire training set.

$$\mathcal{L}(\Theta) = \frac{1}{m} \sum_{i=1}^m L_{CE}(\hat{y}^{(i)}, y^{(i)})$$

Learning

- We'll do this by starting with a random set of weights and then iteratively updating those weights to lower this cost.

$$\mathcal{L}(\Theta) = \frac{1}{m} \sum_{i=1}^m L_{CE}(\hat{y}^{(i)}, y^{(i)})$$

Learning

Updated weights

Current weights

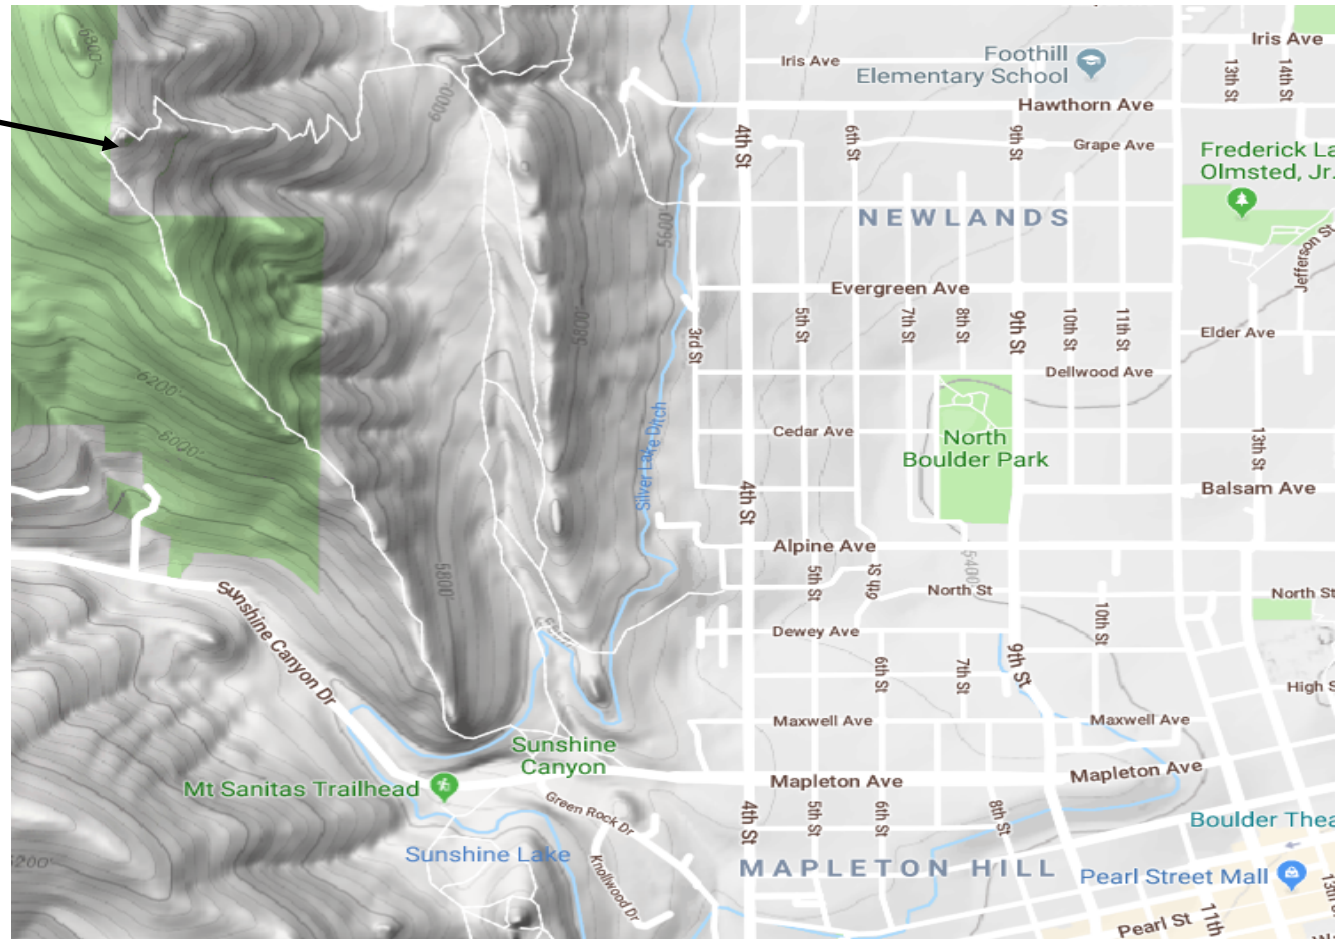
$$w_{t+1} = w_t - \mu \Delta$$

Learning rate

Gradient (vector of partial derivatives of the weights with respect to the loss).

Motivation

This is you. —
Find the fastest
way down.



Derivatives

- Fortunately, we know how to do that from calculus.
- If we take the derivative of the loss function with respect to the weights that will tell us the direction and magnitude of change we should make to each weight.

Derivatives

$$\frac{d}{dx}x^n = nx^{n-1} \quad (1)$$

$$\frac{d}{dx}\cos(x) = -\sin(x) \quad (2)$$

$$\frac{d}{dx}\sin(x) = \cos(x) \quad (3)$$

$$\frac{d}{dx}\log(x) = \frac{1}{x} \quad (4)$$


$$\frac{d}{dx}(b + cx) = c \quad (5)$$

$$\frac{d}{dx}(f(x) + g(x)) = \frac{d}{dx}f(x) + \frac{d}{dx}g(x) \quad (6)$$

$$\frac{d}{dx}(f(x)g(x)) = g(x)\frac{d}{dx}f(x) + f(x)\frac{d}{dx}g(x) \quad (7)$$


Taken from [here](#)

Single Weight



Single Weight

Why do we want to find a minimum?



Partial Derivative

- Of course, in real applications we have many features/weights not just one.
- So we need the vector of the partial derivatives of the loss with respect to the weights.

Cross-Entropy Loss Partial Derivative

$$\frac{\partial L_{CE}(w, b)}{\partial w_j} = [\sigma(w \cdot x + b) - y]x_j$$

Computed answer right answer

Take the difference

Multiply by the feature value of the input corresponding to the weight

Stochastic Gradient Descent

Stochastic Gradient Descent

function STOCHASTIC GRADIENT DESCENT($L()$, $f()$, x , y) **returns** θ

where: L is the loss function

f is a function parameterized by θ

x is the set of training inputs $x^{(1)}, x^{(2)}, \dots, x^{(n)}$

y is the set of training outputs (labels) $y^{(1)}, y^{(2)}, \dots, y^{(n)}$

$\theta \leftarrow 0$

repeat T times

For each training tuple $(x^{(i)}, y^{(i)})$ (in random order)

Compute $\hat{y}^{(i)} = f(x^{(i)}; \theta)$ # What is our estimated output \hat{y} ?

Compute the loss $L(\hat{y}^{(i)}, y^{(i)})$ # How far off is $\hat{y}^{(i)}$ from the true output $y^{(i)}$?

$g \leftarrow \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$ # How should we move θ to maximize loss ?

$\theta \leftarrow \theta - \eta g$ # go the other way instead

return θ

Stochastic Gradient Descent

- Imagine a logistic regression classifier:
 - $w = [2.5, -5, -1.2, 0.5, 2.0, 0.7]$, $b = 0.1$
- Example :
 - $x = [3, 2, 1, 3, 0, 4.15]$, $y = 1$
- Prediction:

Stochastic Gradient Descent

- Imagine a logistic regression classifier:
 - $w = [2.5, -5, -1.2, 0.5, 2.0, 0.7]$, $b = 0.1$
- Example :
 - $x = [3, 2, 1, 3, 0, 4.15]$, $y = 1$
- Prediction:
 - $\text{sigmoid}(2.5*3-2*5-1*1.2+0.5*3+2.0*0+0.7*4.15+0.1) = \text{sigmoid}(0.805) = 0.69$

Stochastic Gradient Descent

- Imagine a logistic regression classifier:
 - $w = [2.5, -5, -1.2, 0.5, 2.0, 0.7]$, $b = 0.1$
- Example :
 - $x = [3, 2, 1, 3, 0, 4.15]$, $y = 1$
- Prediction:
 - $\text{sigmoid}(2.5*3-2*5-1*1.2+0.5*3+2.0*0+0.7*4.15+0.1) = \text{sigmoid}(0.805) = 0.69$
- Loss:

$$L_{CE}(w, b) = -[y \log \sigma(w \cdot x + b) + (1 - y) \log (1 - \sigma(w \cdot x + b))]$$

Stochastic Gradient Descent

- Imagine a logistic regression classifier:
 - $w = [2.5, -5, -1.2, 0.5, 2.0, 0.7]$, $b = 0.1$
- Example :
 - $x = [3, 2, 1, 3, 0, 4.15]$, $y = 1$
- Gradient:
$$\frac{\partial L_{CE}(w, b)}{\partial w_j} = [\sigma(w \cdot x + b) - y]x_j$$
 - $[-0.31 \cdot 3, -0.31 \cdot 2, -0.31 \cdot 1, -0.31 \cdot 3, -0.31 \cdot 0, -0.31 \cdot 4.15] = [-0.93, -0.62, -0.31, -0.93, 0, -1.28]$

Stochastic Gradient Descent

$$w_{t+1} = w_t - \mu \Delta$$

- Current weights:
 - $w = [2.5, -5, -1.2, 0.5, 2.0, 0.7]$
- Gradient:
 - $[-0.93, -0.62, -0.31, -0.93, 0, -1.28]$
- Assume a learning rate of 1 for simplicity!
- New weights:
 - $[3.42687096, -4.38208602, -0.89104301, 1.42687096, 2. , 1.9821715]$

Redo the Example

- New weights:
 - $w = [3.43, -4.38, -0.89, 1.43, 2., 1.98],$
 - $b = 0.1$
- Example (the same as before!):
 - $x = [3, 2, 1, 3, 0, 4.15], y = 1$

Redo the Example

- New weights:
 - $w = [3.43, -4.38, -0.89, 1.43, 2., 1.98],$
 - $b = 0.1$
- Example (the same as before!):
 - $x = [3, 2, 1, 3, 0, 4.15], y = 1$
- Score
 - $\sigma([3.43, -4.38, -0.89, 1.43, 2., 1.98] \cdot [3, 2, 1, 3, 0, 4.15] + 0.1)$
 - $= \sigma(13.247)$
 - $= 1.$

About the Bias

- There's that pesky bias term. That's also a parameter and we forgot to update it as well.

Example

- Imagine a logistic regression classifier:
 - $w = [2.5, -5, -1.2, 0.5, 2.0, 0.7]$, $b = 0.1$
- Example :
 - $x = [3, 2, 1, 3, 0, 4.15]$

Example

- Imagine a logistic regression classifier:
 - $w = [2.5, -5, -1.2, 0.5, 2.0, 0.7, \mathbf{0.1}]$
- Example :
 - $x = [3, 2, 1, 3, 0, 4.15, \mathbf{1}]$

In-Class Exercise 2

- Imagine a logistic regression classifier:
 - $w = [0, 0, 0]$, $b = 0$
- Example 1:
 - $x = [0, 1, 0]$, $y = 0$
- Example 2:
 - $x = [1, 0, 1]$, $y = 1$
- Perform 2 steps of stochastic gradient descent using cross-entropy loss and learning rate 0.1! What parameters do you obtain?

Optimization

- In practice, that can be slow to converge because the algorithm can either be taking steps...
 - ...that are too small and hence take us too long to get where we're going
 - ...or too large which leads us to overshoot the target and wander around too much

Optimization

- In practice, that can be slow to converge because the algorithm can either be taking steps...
 - ...that are too small and hence take us too long to get where we're going
 - ...or too large which leads us to overshoot the target and wander around too much
- Fortunately, you don't have to worry about this. Lots of packages available where you need to specify the loss function and parameters and you're done.

Softmax Loss

- What's the loss for the softmax?
 - We'll use the same cross-entropy idea
 - The probability assigned by the model to the correct class
 - Then use the negative log probability of that

Stochastic Gradient Descent

- Batch training
 - Process each example in the training set and accumulate the gradients
 - Do a single update
 - Repeat

Stochastic Gradient Descent

- Batch training
 - Process each example in the training set and accumulate the gradients
 - Do a single update
 - Repeat
- Minibatch training
 - Select N examples and proceed as with batch training
 - Update after each mini-batch
 - N is chosen to maximize parallelism

Building Your Own Sentiment Classifier

Let's See Some Code!

```
if __name__ == "__main__":
    train, dev = load_data()
    feature_dict = get_features(train)
    train = make_feature_vectors(train, feature_dict)
    dev = make_feature_vectors(dev, feature_dict)

    # TODO: substitute -1 with the correct value!
    model = MyClassifier(len(feature_dict), -1)

    loss_function = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=.1)

    eval(model, train)
    eval(model, dev)
    print()

    for i in range(3):
        model.train()
        for (x, y) in train:
            model.zero_grad()
            raw_scores = model(x)
            loss = loss_function(raw_scores.unsqueeze(0), y)
            loss.backward()
            optimizer.step()

    eval(model, train)
    eval(model, dev)
    print()
```

```
class MyClassifier(nn.Module):

    def __init__(self, num_features, num_labels):
        super(MyClassifier, self).__init__()

        # TODO: substitute -1 with the correct value!
        self.linear = nn.Linear(num_features, -1)

    def forward(self, input):
        return self.linear(input)
```

In-Class Exercise: BYOSC

- Download the sentiment dataset and the prepared code from Canvas (originally from <http://help.sentiment140.com/for-students>)
 - What do the labels mean?
- Fill in the TODOs in the code
 - What is the accuracy of your classifier?
- Improve your classifier!
 - For example: can you find better features?
 - What is the highest accuracy you can obtain?

Wrapping up

- Discussed today:
 - Cross-entropy loss
 - Stochastic gradient descent
- On Monday: Training neural networks, part 2